

# 50.012 Networks

## Lecture 4: CDN and P2P

2021 Term 6

Assoc. Prof. CHEN Binbin



# Learning objectives

Understand the design of content distribution networks (CDN), including the design objectives, key design considerations, and various approaches.

Understand the peer-to-peer architecture: benefits, challenges, basic operations and incentive mechanism of BitTorrent.

Read: textbook Section 2.4, 2.5, 2.6

# Content distribution networks


- **Challenge:** how to stream content (selected from millions of videos) to hundreds of thousands of *simultaneous* users?
    - Recall: they're highly heterogeneous, including locations, connectivity, etc.
  - **Option 1:** single, large “mega-server”
    - point of network congestion
    - long path to distant clients
    - multiple copies of video sent over outgoing link (wastage)
    - single point of failure
- ....quite simply: this solution *doesn't scale*

# Content distribution networks

- *Challenge:* how to stream content (selected from millions of videos) to hundreds of thousands of simultaneous users?
- *Option 2:* store/serve multiple copies of videos at multiple **geographically distributed** sites (*CDN*)

# Content distribution networks

- CDN: stores copies of content at CDN nodes
  - e.g., Netflix stores copies of MadMen
- subscriber requests content from Netflix
  - directed to nearby CDN copy (based on mapping IP address of request to geographical location), retrieves content

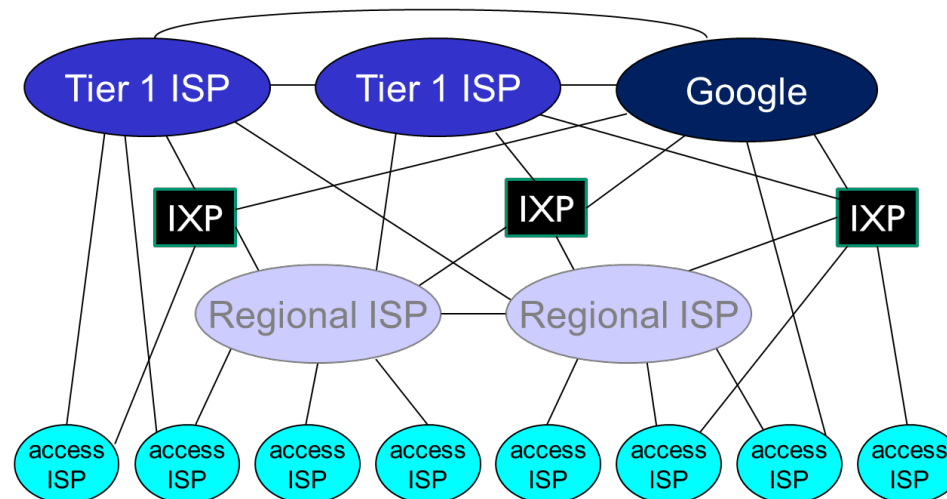


# Questions for a CDN operator

- Where to put the servers?
- How to select the server?
- How to route the client to the server?
- How to replicate the content to the servers?

# CDN Server Placement

- *Enter deep*: push CDN servers deep into many access ISPs
  - close to users
  - used by Akamai, 1700+ locations
- *Bring home*: smaller number (10's) of larger clusters in IXPs (Internet Exchange Points) near (but not within) access networks
  - used by Limelight



# Types of CDN

- Commercial CDN
  - Akamai Technologies, Cloudflare, LimeLight
- Content provider's own CDN
  - Google, netflix
- Telco CDN



# Cluster (Server) Selection

- Geographically close
- Best performance: real-time measurement
- Lowest load: Load-balancing
- Cheapest: CDN may need to pay its provider ISP
- Any alive node: fault-tolerant

# Content Access / Routing

- Naming (DNS)–based
- Application-driven
  - E.g., use HTTP redirect
  - Multiple connection setup, name lookups
- Routing (anycast)-based
  - Coarse-grained


# A quick recap of DNS (Domain Name System)

# DNS name resolution example

- host at cis.poly.edu wants IP address for gaia.cs.umass.edu

## *iterated query:*

- contacted server replies with name of server to contact
- “I don’t know this name, but ask this server”



# DNS records

**DNS:** distributed database storing resource records (**RR**)

RR format: (**name**, **value**, **type**, **ttl**)

## type=A

- **name** is hostname
- **value** is IP address

## type=NS

- **name** is domain (e.g., foo.com)
- **value** is hostname of authoritative name server for this domain

## type=CNAME

- **name** is alias name for some “canonical” (the real) name
- **www.ibm.com** is really **servereast.backup2.ibm.com**
- **value** is canonical name

## type=MX


- **value** is name of mailserver associated with **name**

# CDN content access by DNS redirect

Bob (client) requests video `http://video.netcinema.com/6Y7B23V`

- video stored in CDN run by KingCDN.com

Used by: third-party CDNs like Akamai, if origin URLs aren't "Akamaized"



# CDN content access by DNS redirect

Bob (client) requests video `http://video.netcinema.com/6Y7B23V`

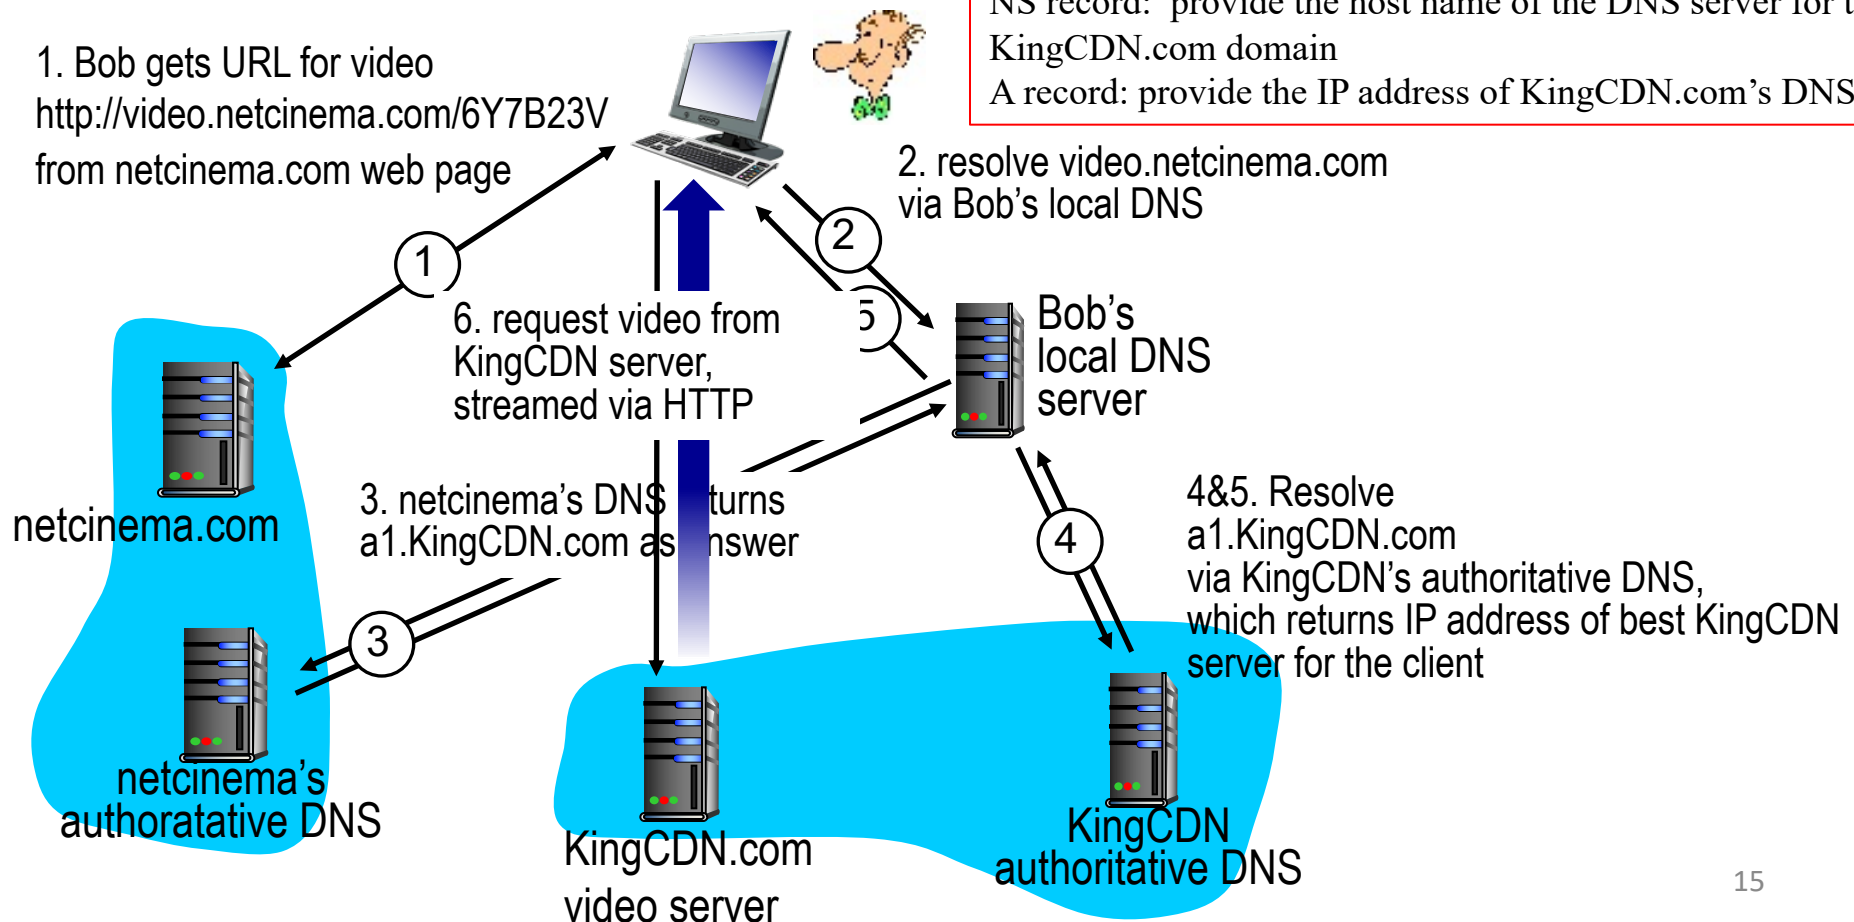
- video stored in CDN run by KingCDN.com

The DNS records sent at step 3:

CNAME record: map `video.netcinema.com` to `a1.KingCDN.com`

NS record: provide the host name of the DNS server for the KingCDN.com domain

A record: provide the IP address of KingCDN.com's DNS server



# Akamai Resource Locators (ARL)

- Each customer has its dedicated domain name
  - E.g., a123.g.akamaitech.net
  - Akamai operates the authoritative DNS servers for these names
- The original server can “akamaize” the origin URL to ARL (Akamai Resource Locators)
  - E.g.,: `` becomes ``
- Client browser issues GET to CDN instead of origin server
- Read more:  
Akamai “Fast Internet Content Delivery with FreeFlow”




# Case study: Netflix

- Netflix used to rely on third-party CDNs
  - Akamai, LimeLight, level-3
- It has built its private CDN for delivering videos
  - Still rely on Akamai for delivering web pages
  - Server placements: bring home (50+ IXP) + enter deep (free rack of servers for access ISPs)
- Netflix rack: 1 (access ISP) to 10+ (IXP) servers
  - Each server has multiple 10Gbps Ethernet ports, 100 Terabytes+ storage

# Case study: Netflix

NB. CDN is owned (hence directly controlled) by Netflix: no need for previous DNS redirect



# Content replication

	Mechanism	Pros.	Cons.
Netflix	Push		
Youtube	Pull		

Why the difference in their design?

# Use of CDN

- DDoS Attack mitigation
- Handle flash crowds
- Web application firewalls (WAF)

# Outline

Content distribution networks (CDN)


Peer-to-peer architecture

# Pure P2P architecture

- No always-on server
- Arbitrary end systems (e.g., our laptops) directly communicate as peers
- Peers are intermittently connected and may change IP addresses

## Examples:


- File (e.g., video) distribution (BitTorrent)
- Multimedia streaming (KanKan )
- VoIP (Skype)



# File distribution: client-server vs P2P

Question: How much time to distribute file (size  $F$ ) from one server to  $N$  peers?

- peer upload/download capacity (access bandwidth) is limited resource, compared with core bandwidth




# File distribution time (lower bound): client-server

- **server transmission:** must sequentially send (upload)  $N$  file copies:

- time to send one copy:  $F/u_s$
- time to send  $N$  copies:  $NF/u_s$

- **client:** each client must download file copy

- $d_{min}$  = min client download rate
- min client download time:  $F/d_{min}$



*time to distribute  $F$   
to  $N$  clients using  
client-server approach*


$$D_{c-s} \geq \max\{NF/u_s, F/d_{min}\}$$

increases linearly in  $N$ ;  
susceptible to “flash  
crowds” particularly



# File distribution time (lower bound): P2P

- **server transmission:** must upload at least one copy
  - time to send one copy:  $F/u_s$
- **client:** each client must download a file copy
  - min client download time:  $F/d_{\min}$
- **clients:** as aggregate must download  $NF$  bits
  - implies upload of the same number of bits (**why?**)
  - max upload rate is  $u_s + \sum u_i$  (if it's feasible to schedule the distribution to make all the uploads concurrently active)



time to distribute  $F$   
to  $N$  clients using  
P2P approach

$$D_{P2P} \geq \max\{F/u_s, F/d_{\min}, NF/(u_s + \sum u_i)\}$$

increases linearly in  $N$  ...


... but so does this, as each peer brings service capacity

# P2P file distribution: BitTorrent

- file divided into 256 kB chunks
- peers in torrent send/receive (trade) file chunks


*tracker*: tracks peers participating in torrent

*torrent*: group of peers exchanging chunks of a file




# Client-server vs. P2P: example

client upload rate =  $u$ ,  $F/u = 1$  hour,  $u_s = 10u$ ,  $d_{min} \geq u_s$



# P2P file distribution: BitTorrent

- Peer joining torrent:
  - initially has no chunks, but will accumulate them over time from other peers
  - registers with tracker to get list of peers, connects to subset of peers (“neighbors”)



- While downloading, peer uploads chunks to other peers
- Peer may change peers with whom it exchanges chunks
- **Churn:** peers may come and go (highly dynamic & organic, *no a priori* set up reliable infrastructure)
- Once peer has entire file, it may (selfishly) leave or (altruistically) remain in torrent
  - peers ultimately interested in download
  - In general need *incentives* for them to provide upload (e.g., tit-for-tat)

# BitTorrent: requesting, sending file chunks

## Requesting chunks:

- at any given time, different peers have different subsets of file chunks
- periodically, Alice asks each peer for list of chunks that they have
- Alice requests missing chunks from peers, **rarest first**
  - Aim to make all the chunks equally available (they are equally important)

## Sending chunks: *tit-for-tat*

- Alice sends chunks to those four peers currently sending her chunks *at highest rate*
  - these 4 are “unchoked”
  - other peers are choked by Alice (do not receive chunks from her)
  - re-evaluate top 4 every 10 secs
- every 30 secs: randomly select another peer, starts sending chunks
  - optimistically “unchoke” this peer: give chance of discovering an even better trading partner than the original top 4
  - newly chosen peer may join top 4

# BitTorrent: tit-for-tat

- (1) Alice “optimistically unchokes” Bob
- (2) Alice becomes one of Bob’s top-four providers; Bob reciprocates
- (3) Bob in turn becomes one of Alice’s top-four providers

